

```
In [1]: import os
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: os.getcwd()
```

```
Out[2]: 'C:\\Users\\Preksha\\Simplilearn\\Data Science with Python\\MovieLens'
```

```
In [3]: movie = ['MovieID', 'Title', 'Genres']
```

```
In [4]: rating = ['UserID', 'MovieID', 'Rating', 'Timestamp']
```

```
In [5]: user = ['UserID', 'Gender', 'Age', 'Occupation', 'Zip-code']
```

Import the three datasets

```
In [6]: movies = pd.read_csv('movies.dat', delimiter = '::', names = movie)
```

<ipython-input-6-7f97736ab933>:1: ParserWarning: Falling back to the 'python' engine because the 'c' engine does not support regex separators (separators > 1 char and different from '\s+' are interpreted as regex); you can avoid this warning by specifying engine='python'.

```
movies = pd.read_csv('movies.dat', delimiter = '::', names = movie)
```

```
In [7]: ratings = pd.read_csv('ratings.csv', names = rating)
```

```
In [8]: users = pd.read_csv('users.csv', names = user)
```

```
In [9]: movies.head()
```

```
Out[9]:
```

	MovieID	Title	Genres
0	1	Toy Story (1995)	Animation Children's Comedy
1	2	Jumanji (1995)	Adventure Children's Fantasy
2	3	Grumpier Old Men (1995)	Comedy Romance
3	4	Waiting to Exhale (1995)	Comedy Drama
4	5	Father of the Bride Part II (1995)	Comedy

In [10]: `ratings.head()`

Out[10]:

	UserID	MovieID	Rating	Timestamp
0	1	1193	5	978300760
1	1	661	3	978302109
2	1	914	3	978301968
3	1	3408	4	978300275
4	1	2355	5	978824291

In [11]: `users.head()`

Out[11]:

	UserID	Gender	Age	Occupation	Zip-code
0	1	F	1	10	48067
1	2	M	56	16	70072
2	3	M	25	15	55117
3	4	M	45	7	02460
4	5	M	25	20	55455

Create a new dataset [Master_Data] with the following columns MovieID Title UserID Age Gender Occupation Rating. (Hint: (i) Merge two tables at a time. (ii) Merge the tables using two primary keys MovieID & UserID)

In [12]: `master_data = pd.merge(ratings, movies, how='outer', on='MovieID')`

In [13]: `master_data.head()`

Out[13]:

	UserID	MovieID	Rating	Timestamp	Title	Genres
0	1.0	1193	5.0	978300760.0	One Flew Over the Cuckoo's Nest (1975)	Drama
1	2.0	1193	5.0	978298413.0	One Flew Over the Cuckoo's Nest (1975)	Drama
2	12.0	1193	4.0	978220179.0	One Flew Over the Cuckoo's Nest (1975)	Drama
3	15.0	1193	4.0	978199279.0	One Flew Over the Cuckoo's Nest (1975)	Drama
4	17.0	1193	5.0	978158471.0	One Flew Over the Cuckoo's Nest (1975)	Drama

```
In [14]: master_data1 = pd.merge(master_data, users, how = 'outer', on = 'UserID')
```

```
In [15]: master_data1.head()
```

Out[15]:

	UserID	MovieID	Rating	Timestamp	Title	Genres	Gender	Age
0	1.0	1193	5.0	978300760.0	One Flew Over the Cuckoo's Nest (1975)	Drama	F	1.0
1	1.0	661	3.0	978302109.0	James and the Giant Peach (1996)	Animation Children's Musical	F	1.0
2	1.0	914	3.0	978301968.0	My Fair Lady (1964)	Musical Romance	F	1.0
3	1.0	3408	4.0	978300275.0	Erin Brockovich (2000)	Drama	F	1.0
4	1.0	2355	5.0	978824291.0	Bug's Life, A (1998)	Animation Children's Comedy	F	1.0

```
In [16]: master_data = master_data1.drop(['Timestamp', 'Genres', 'Zip-code'], axis = 1)
```

```
In [17]: master_data.head()
```

Out[17]:

	UserID	MovieID	Rating	Title	Gender	Age	Occupation
0	1.0	1193	5.0	One Flew Over the Cuckoo's Nest (1975)	F	1.0	10.0
1	1.0	661	3.0	James and the Giant Peach (1996)	F	1.0	10.0
2	1.0	914	3.0	My Fair Lady (1964)	F	1.0	10.0
3	1.0	3408	4.0	Erin Brockovich (2000)	F	1.0	10.0
4	1.0	2355	5.0	Bug's Life, A (1998)	F	1.0	10.0

```
In [18]: master_data.isna().sum()
```

```
Out[18]: UserID      177
MovieID      0
Rating      177
Title        0
Gender      177
Age         177
Occupation   177
dtype: int64
```

```
In [19]: master_data.size
```

```
Out[19]: 7002702
```

```
In [20]: master_data.shape
```

```
Out[20]: (1000386, 7)
```

```
In [21]: master_data.columns
```

```
Out[21]: Index(['UserID', 'MovieID', 'Rating', 'Title', 'Gender', 'Age', 'Occupation'], dtype='object')
```

```
In [22]: master_data.dtypes
```

```
Out[22]: UserID      float64
MovieID      int64
Rating      float64
Title       object
Gender      object
Age        float64
Occupation  float64
dtype: object
```

```
In [23]: master_data.dropna(inplace = True)
```

```
In [24]: master_data.isna().sum()
```

```
Out[24]: UserID      0
MovieID      0
Rating      0
Title      0
Gender      0
Age      0
Occupation  0
dtype: int64
```

```
In [25]: master_data.corr()
```

```
Out[25]:
```

	UserID	MovieID	Rating	Age	Occupation
UserID	1.000000	-0.017739	0.012303	0.034688	-0.026698
MovieID	-0.017739	1.000000	-0.064042	0.027575	0.008585
Rating	0.012303	-0.064042	1.000000	0.056869	0.006753
Age	0.034688	0.027575	0.056869	1.000000	0.078371
Occupation	-0.026698	0.008585	0.006753	0.078371	1.000000


```
In [26]: master_data.nunique()
```

```
Out[26]: UserID      6040  
MovieID      3706  
Rating        5  
Title      3706  
Gender        2  
Age           7  
Occupation    21  
dtype: int64
```

Explore the datasets using visual representations (graphs or tables), also include your comments on the following:

User Age Distribution

```
In [27]: master_data.Age.hist(grid = False);
```



```
In [28]: rating_movies= master_data.drop(['UserID', 'MovieID', 'Gender', 'Age', 'Occupation'])
```

```
In [29]: rating_movies.head()
```

```
Out[29]:
```

	Rating	Title
0	5.0	One Flew Over the Cuckoo's Nest (1975)
1	3.0	James and the Giant Peach (1996)
2	3.0	My Fair Lady (1964)
3	4.0	Erin Brockovich (2000)
4	5.0	Bug's Life, A (1998)

```
In [30]: toy_story = rating_movies[rating_movies['Title'].str.startswith('Toy Story'),
toy_story
```


Out[30]:

	Rating	Title
40	5.0	Toy Story (1995)
50	4.0	Toy Story 2 (1999)
417	5.0	Toy Story 2 (1999)
634	4.0	Toy Story (1995)
938	5.0	Toy Story (1995)
...
994256	5.0	Toy Story 2 (1999)
994289	5.0	Toy Story 2 (1999)
994315	4.0	Toy Story 2 (1999)
994367	4.0	Toy Story 2 (1999)
994389	4.0	Toy Story 2 (1999)

3662 rows × 2 columns

User rating of the movie “Toy Story”

```
In [31]: toy_story.Rating.hist(grid=False);
```



Top 25 movies by viewership rating

```
In [32]: top_25 = rating_movies.sort_values('Rating', ascending = False)
top_25 = top_25.head(25)
```

In [33]: `top_25`

Out[33]:

	Rating	Title
0	5.0	One Flew Over the Cuckoo's Nest (1975)
327573	5.0	Crucible, The (1996)
327567	5.0	Blow-Out (La Grande Bouffe) (1973)
327564	5.0	Five Easy Pieces (1970)
830521	5.0	Black Cauldron, The (1985)
327559	5.0	Raging Bull (1980)
327554	5.0	Cook the Thief His Wife & Her Lover, The (1989)
830530	5.0	Great Escape, The (1963)
830539	5.0	Conan the Barbarian (1982)
830543	5.0	Sneakers (1992)
830547	5.0	Devil's Advocate, The (1997)
830552	5.0	Space Cowboys (2000)
327507	5.0	Conquest of the Planet of the Apes (1972)
830561	5.0	Secret of NIMH, The (1982)
327490	5.0	Manhattan Murder Mystery (1993)
327487	5.0	Love and Death (1975)
327479	5.0	Mother (1996)
830578	5.0	On Her Majesty's Secret Service (1969)
830583	5.0	Fox and the Hound, The (1981)
830520	5.0	Tron (1982)
327582	5.0	Lolita (1962)
327306	5.0	Jungle Fever (1991)
830514	5.0	Blade Runner (1982)
830468	5.0	Fantasia (1940)
830469	5.0	Dr. Strangelove or: How I Learned to Stop Worr...

In [34]: `master_data.columns`

Out[34]: Index(['UserID', 'MovieID', 'Rating', 'Title', 'Gender', 'Age', 'Occupation'], dtype='object')

In [35]: `users_rating = master_data.drop(['MovieID', 'Gender', 'Age', 'Occupation'], ax`

```
In [36]: users_rating.head()
```

Out[36]:

	UserID	Rating	Title
0	1.0	5.0	One Flew Over the Cuckoo's Nest (1975)
1	1.0	3.0	James and the Giant Peach (1996)
2	1.0	3.0	My Fair Lady (1964)
3	1.0	4.0	Erin Brockovich (2000)
4	1.0	5.0	Bug's Life, A (1998)

**Find the ratings for all the movies reviewed by
for a particular user of user id = 2696**

```
In [37]: users_rating = users_rating[users_rating['UserID'] == 2696 ]
```




```
In [38]: users_rating.sort_values('Rating', ascending = False)
```

Out[38]:

	UserID	Rating	Title
953850	2696.0	5.0	Lone Star (1996)
953857	2696.0	4.0	Game, The (1997)
953859	2696.0	4.0	Devil's Advocate, The (1997)
953849	2696.0	4.0	L.A. Confidential (1997)
953864	2696.0	4.0	Shining, The (1980)
953852	2696.0	4.0	Talented Mr. Ripley, The (1999)
953853	2696.0	4.0	Midnight in the Garden of Good and Evil (1997)
953862	2696.0	4.0	Basic Instinct (1992)
953855	2696.0	4.0	Palmetto (1998)
953856	2696.0	4.0	Perfect Murder, A (1998)
953861	2696.0	4.0	Wild Things (1998)
953860	2696.0	4.0	Psycho (1998)
953866	2696.0	3.0	Client, The (1994)
953848	2696.0	3.0	E.T. the Extra-Terrestrial (1982)
953854	2696.0	3.0	Cop Land (1997)
953858	2696.0	2.0	I Know What You Did Last Summer (1997)
953865	2696.0	2.0	I Still Know What You Did Last Summer (1998)
953847	2696.0	2.0	Back to the Future (1985)
953863	2696.0	1.0	Lake Placid (1999)
953851	2696.0	1.0	JFK (1991)

```
In [39]: users_rating.Rating.hist();
```



```
In [40]: master_data1.Genres.unique()
```

```
Out[40]: array(['Drama', 'Animation|Children's|Musical', 'Musical|Romance',  
               'Animation|Children's|Comedy', 'Action|Adventure|Comedy|Romance',  
               'Action|Adventure|Drama', 'Comedy|Drama',  
               'Adventure|Children's|Drama|Musical', 'Musical', 'Comedy',  
               'Animation|Children's', 'Comedy|Fantasy', 'Animation',  
               'Comedy|Sci-Fi', 'Drama|War', 'Romance',  
               'Animation|Children's|Musical|Romance',  
               'Children's|Drama|Fantasy|Sci-Fi', 'Drama|Romance',  
               'Animation|Comedy|Thriller',  
               'Adventure|Animation|Children's|Comedy|Musical',  
               'Animation|Children's|Comedy|Musical', 'Thriller',  
               'Action|Crime|Romance', 'Action|Adventure|Fantasy|Sci-Fi',  
               'Children's|Comedy|Musical', 'Action|Drama|War',  
               'Children's|Drama', 'Crime|Drama|Thriller', 'Action|Crime|Drama',  
               'Action|Adventure|Mystery', 'Crime|Drama',  
               'Action|Adventure|Sci-Fi|Thriller',  
               'Action|Adventure|Romance|Sci-Fi|War', 'Action|Thriller',  
               'Action|Drama', 'Comedy|Drama|Western', 'Action|Adventure|Crime',  
               'Action|Crime|Mystery|Thriller', 'Comedy|Drama|Romance',  
               ...])
```

```
In [41]: master_data1.isna().sum()
```

```
Out[41]: UserID      177  
         MovieID      0  
         Rating      177  
         Timestamp    177  
         Title        0  
         Genres        0  
         Gender      177  
         Age          177  
         Occupation    177  
         Zip-code     177  
         dtype: int64
```

```
In [42]: master_data1.dropna(inplace = True)
```

```
In [43]: master_data1.isna().any()
```

```
Out[43]: UserID      False  
         MovieID      False  
         Rating      False  
         Timestamp    False  
         Title        False  
         Genres        False  
         Gender        False  
         Age          False  
         Occupation    False  
         Zip-code      False  
         dtype: bool
```

In [44]: `master_data1.head()`

Out[44]:

	UserID	MovieID	Rating	Timestamp	Title	Genres	Gender	Age
0	1.0	1193	5.0	978300760.0	One Flew Over the Cuckoo's Nest (1975)	Drama	F	1.0
1	1.0	661	3.0	978302109.0	James and the Giant Peach (1996)	Animation Children's Musical	F	1.0
2	1.0	914	3.0	978301968.0	My Fair Lady (1964)	Musical Romance	F	1.0
3	1.0	3408	4.0	978300275.0	Erin Brockovich (2000)	Drama	F	1.0
4	1.0	2355	5.0	978824291.0	Bug's Life, A (1998)	Animation Children's Comedy	F	1.0

Find out all the unique genres (Hint: split the data in column genre making a list and then process the data to find out only the unique categories of genres)

In [45]: `Genres = master_data1['Genres'].str.split('|')`

In [46]: `list1 = Genres.tolist()`


```
In [51]: master_data2 = pd.concat([master_data1, master_data1.Genres.str.get_dummies()])
print(master_data2.columns)
```


```
Index(['UserID', 'MovieID', 'Rating', 'Timestamp', 'Title', 'Genres', 'Gender',
      'Age', 'Occupation', 'Zip-code', 'Action', 'Adventure', 'Animation',
      'Children's', 'Comedy', 'Crime', 'Documentary', 'Drama', 'Fantasy',
      'Film-Noir', 'Horror', 'Musical', 'Mystery', 'Romance', 'Sci-Fi',
      'Thriller', 'War', 'Western'],
      dtype='object')
```

```
In [52]: master_data2.head()
```

Out[52]:

	UserID	MovieID	Rating	Timestamp	Title	Genres	Gender	Age
0	1.0	1193	5.0	978300760.0	One Flew Over the Cuckoo's Nest (1975)	Drama	F	1.0
1	1.0	661	3.0	978302109.0	James and the Giant Peach (1996)	Animation Children's Musical	F	1.0
2	1.0	914	3.0	978301968.0	My Fair Lady (1964)	Musical Romance	F	1.0
3	1.0	3408	4.0	978300275.0	Erin Brockovich (2000)	Drama	F	1.0
4	1.0	2355	5.0	978824291.0	Bug's Life, A (1998)	Animation Children's Comedy	F	1.0

5 rows × 28 columns



```
In [53]: master_data2.Occupation.unique()
```

Out[53]: array([10., 16., 12., 7., 1., 3., 4., 8., 17., 0., 2., 9., 19., 18., 15., 11., 20., 13., 5., 14., 6.])

Create a separate column for each genre category with a one-hot encoding (1 and 0) whether or not the movie belongs to that genre.

```
In [54]: master_data2 = master_data2.drop(['MovieID', 'UserID', 'Timestamp', 'Zip-code'])
```

```
In [55]: master_data2 = pd.get_dummies(master_data2, columns=['Gender', 'Occupation'])
```

```
In [56]: master_data2
```

Out[56]:

	Rating	Age	Action	Adventure	Animation	Children's	Comedy	Crime	Documentar
0	5.0	1.0	0	0	0	0	0	0	
1	3.0	1.0	0	0	1	1	0	0	
2	3.0	1.0	0	0	0	0	0	0	
3	4.0	1.0	0	0	0	0	0	0	
4	5.0	1.0	0	0	1	1	1	0	
...	
1000204	2.0	45.0	0	0	0	0	0	0	
1000205	3.0	45.0	0	0	0	0	0	0	
1000206	4.0	45.0	0	0	0	0	0	0	
1000207	2.0	45.0	1	0	0	0	0	0	
1000208	2.0	45.0	0	1	0	0	0	0	

1000209 rows × 43 columns

Determine the features affecting the ratings of any particular movie.

```
In [57]: master_data2.columns
```

Out[57]: Index(['Rating', 'Age', 'Action', 'Adventure', 'Animation', 'Children's', 'Comedy', 'Crime', 'Documentary', 'Drama', 'Fantasy', 'Film-Noir', 'Horror', 'Musical', 'Mystery', 'Romance', 'Sci-Fi', 'Thriller', 'Western', 'Gender_F', 'Gender_M', 'Occupation_0.0', 'Occupation_1.0', 'Occupation_2.0', 'Occupation_3.0', 'Occupation_4.0', 'Occupation_5.0', 'Occupation_6.0', 'Occupation_7.0', 'Occupation_8.0', 'Occupation_9.0', 'Occupation_10.0', 'Occupation_11.0', 'Occupation_12.0', 'Occupation_13.0', 'Occupation_14.0', 'Occupation_15.0', 'Occupation_16.0', 'Occupation_17.0', 'Occupation_18.0', 'Occupation_19.0', 'Occupation_20.0'], dtype='object')

```
In [58]: from sklearn.model_selection import train_test_split
```

```
In [59]: X = master_data2[['Age', 'Action', 'Adventure', 'Animation',  
    "Children's", 'Comedy', 'Crime', 'Documentary', 'Drama', 'Fantasy',  
    'Film-Noir', 'Horror', 'Musical', 'Mystery', 'Romance', 'Sci-Fi',  
    'Thriller', 'War', 'Western', 'Gender_F', 'Gender_M', 'Occupation_0.0',  
    'Occupation_1.0', 'Occupation_2.0', 'Occupation_3.0', 'Occupation_4.0',  
    'Occupation_5.0', 'Occupation_6.0', 'Occupation_7.0', 'Occupation_8.0',  
    'Occupation_9.0', 'Occupation_10.0', 'Occupation_11.0',  
    'Occupation_12.0', 'Occupation_13.0', 'Occupation_14.0',  
    'Occupation_15.0', 'Occupation_16.0', 'Occupation_17.0',  
    'Occupation_18.0', 'Occupation_19.0', 'Occupation_20.0']]
```

```
In [60]: y = master_data2[['Rating']]
```

```
In [61]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.75,
```

```
In [62]: X_train.shape, X_test.shape, y_train.shape, y_test.shape
```

```
Out[62]: ((750156, 42), (250053, 42), (750156, 1), (250053, 1))
```

Develop an appropriate model to predict the movie ratings

```
In [63]: from sklearn.tree import DecisionTreeClassifier
```

```
In [64]: dtc_model = DecisionTreeClassifier(criterion = 'entropy', random_state = 0)
```

```
In [65]: dtc_model.fit(X_train, y_train)
```

```
Out[65]: DecisionTreeClassifier(criterion='entropy', random_state=0)
```

```
In [66]: y_pred = dtc_model.predict(X_test)
```

```
In [67]: from sklearn import metrics
```

```
In [68]: np.sqrt(metrics.mean_squared_error(y_test, y_pred))
```

```
Out[68]: 1.2361860826434847
```